

Homework 3 — Due May 9 11PM

This assignment builds on the previous assignment's theme of examining memory access patterns. You will implement a finite difference solver for the 2D heat diffusion equation in different ways to examine the performance characteristics of the different implementations.

Background on the heat diffusion PDE. The heat diffusion PDE that we will be solving can be written:

$$\frac{\partial T}{\partial t} = \frac{\partial^2 T}{\partial x^2} + \frac{\partial^2 T}{\partial y^2}$$

To solve this PDE, we are going to discretize both the temporal and the spatial derivatives. To do this, we define a two-dimensional¹ grid $G_{i,j}$, $1 \leq i \leq n_x$, $1 \leq j \leq n_y$, where we denote by n_x (resp. n_y) the number of points on the x -axis (resp. y -axis). At each time-step, we will evaluate the temperature and its derivatives at these gridpoints.

While we will consistently use a first order discretization scheme for the temporal derivative, we will use 2nd, 4th or 8th order discretization of the spatial derivative.²

If we denote by $T_{i,j}^t$ the temperature at time t at point (i, j) of the grid, the 2nd order spatial discretization scheme can be written as:

$$T_{i,j}^{t+1} = T_{i,j}^t + C^{(x)} (T_{i+1,j}^t - 2T_{i,j}^t + T_{i-1,j}^t) + C^{(y)} (T_{i,j+1}^t - 2T_{i,j}^t + T_{i,j-1}^t)$$

The $C^{(x)}$ (`xfcl` in the code) and $C^{(y)}$ (`yfcl` in the code) constants are called Courant numbers. They depend on the temporal discretization step, as well as the spatial discretization step. To ensure the stability of the discretization scheme, they have to be less than a maximum value given by the Courant-Friedrichs-Lewy condition.³ You do not have to worry about this, because the starter code already takes care of picking the temporal discretization step as to maximize the Courant numbers while ensuring stability.

The starter code also contains host and device functions named `stencil` which contain the coefficients that go into the update equation. Therefore, you do not need to figure out how to implement the different order updates. You only need to understand how this function works and pass in the arguments correctly.

Boundary conditions. The starter code contains the functions that will update the boundary conditions for you (see file `BC.h`, in particular the function `gpuUpdateBCsOnly`) and the points that are in the border (which has a size of `order / 2`). This way, you do not have to worry about the size of the stencil as you approach the wall.⁴

Various implementations. In this programming assignment, we are going to implement 2 different kernels (and you can do a third one for extra credit):

- Global (function `gpuStencil`): this kernel will use global memory and each thread will update exactly one point of the mesh. You should use a 2D grid with $n_x \times n_y$ threads total.
- Block (function `gpuStencilLoop`): this kernel will also use global memory. Each thread will update `numYPerStep` points of the mesh (these points form a vertical line). You should also use a 2D grid with $n_x \times n_y / \text{numYPerStep}$ threads total.
- (Extra Credit) Shared (function `gpuShared`): this kernel will use shared memory. A group of threads must load a piece of the mesh in shared memory and then compute the new values of the temperatures on the mesh.⁵ Each thread will load and update several elements.

¹In fact, the size of the grid is $g_x \times g_y$ but we will only update the interior region, which size is $n_x \times n_y$.

²For instance, if we use a 4th order scheme, we will express the derivative with respect to x of T at point (i, j) using $T_{i-2,j}^t$, $T_{i-1,j}^t$, $T_{i,j}^t$, $T_{i+1,j}^t$ and $T_{i+2,j}^t$.

³If you are interested, you can read more about this at http://en.wikipedia.org/wiki/Courant-Friedrichs-Lewy_condition.

⁴A general way of dealing with this problem is to change the size of the stencil as you approach the wall. This is complicated and we simplified the process for this homework.

⁵Note however, that the threads that loaded data on the borders of the small piece will stay idle during the computation step.

Parameter file. The parameters used in the computation are read from the file `params.in`. You will need to modify *some parameters* (see description of the starter code) in this file to answer some of the questions. But this file will not be submitted through the submission script.

Here is a list of parameters that are used in the order they appear in the file:

```
int nx_, ny_;    // number of grid points in each dimension
double lx_, ly_; // extent of physical domain in each dimension
int iters_;      // number of iterations to do
int order_;     // order of discretization
```

Starter code The starter code is composed of the following files:

- `main.cu` — This is the CUDA driver file. Do not modify this file.
- `gpuStencil.cu` — This is the file containing the kernels. You will need to modify this file.
- `Makefile` — `make` will build the binary. `make clean` will remove build files as well as debug output. Do not modify this file.
- `params.in` — This file contains a basic set of parameters. For debugging, performance testing, and to answer the questions, you will need to modify this file. The only parameters you should modify are `nx`, `ny` (line 1) and `order` (line 4). This file however will not get submitted through the script.
- `simParams.h` and `simParams.cpp` — These files contain the data structure necessary to handle the parameters of the problem. Do not modify these files.
- `Grid.h` and `Grid.cu` — These files contain the data structure that models the grid used to solve the PDE. Do not modify this file.
- `BC.h` — This file contains the class `boundary_conditions` that will allow us to update the boundary conditions during the simulation. Do not modify this file.
- `hw3.sh` — This script is used to submit jobs to the queue (see Appendix B).

Note The files in the starter code contain some additional information about the implementation in the form of comments. Additionally, the CPU implementation should provide a clearer picture of the method and should aid your GPU work.

Running the program Type `make` to compile the code. Once this is done, you can use the command:

```
$ ./main [-g] [-b] [-s]
```

where `-g` stands for global, `-b` for block, and `-s` for shared. Note that you can run several implementations at the same time (for instance `./main -g -b` to run the global and block implementations).

The output produced by the program will contain:

- The time and bandwidth for the CPU implementation, as well as for the implementations that you specified when running the program.
- A report of the errors for the GPU implementations. Namely, the program will output:
 - The L_2 norm of the final solution from the CPU implementation (*i.e.* the reference).
 - The L_∞ norm of the relative error between the reference solution and the GPU solution.
 - The L_2 norm of the error normalized by the L_2 norm of the CPU implementation.

Typical ranges for the errors (and for the parameters that you will be using) are:

- $[10^{-7}, 10^{-5}]$ for L_∞ norm error.
- $[10^{-8}, 10^{-6}]$ for L_2 norm error.

Search for `TODO` in `gpuStencil.cu` to see where you need to implement code.

Question 1

(30 points) Implement the function `gpuStencil` that runs the solver using global memory. You must also fill in the function `gpuComputation`. The difference (in terms of the norms) between your implementation and the reference to should be in the expected range.

Question 2

(35 points) Implement the function `gpuStencilLoop` that runs the solver using global memory but where each thread computes several points on the grid. You must also fill in the function `gpuComputationLoop`. The difference (in terms of the norms) between your implementation and the reference to should be in the expected range.

Question 3

(15 points) Plot the bandwidth (GB/s) as a function of the grid size (in units of MegaPoints) for the following grid sizes: 256×256 ; 512×512 ; 1024×1024 ; 2048×2048 ; 4096×4096 .

You must have 2 plots (or 3 plots if you choose to do the extra credit) as follows:

1. For `order = 4`, plot the bandwidth for the 2 (or 3) different algorithms.
2. For the block implementation, plot the bandwidth for the 3 different orders.
3. If you implemented the shared algorithm, plot the bandwidth for the 3 different orders.

Question 4

(20 points) Which kernel (global, block or shared) and which order gives the best performance (your answer may depend on the grid size)? Explain the performance results you got in Question 3.

Question 5

(20 points Extra Credit) Implement the function `gpuShared` that runs the solver using shared memory. You should also fill in the function `gpuComputationShared`. Note that you have to answer the questions related to shared memory implementation in Questions 3 and 4 in order to get the full extra credit.

Total number of points: 100 (+ 20 extra credit)

A Submission instructions

To submit:

1. For all questions that require explanations and answers besides source code, put those explanations and answers in a separate PDF file. The name of the file should be: `hw3.pdf`. Please upload the PDF file to Gradescope, **not** `cardinal.stanford.edu`.
2. The homework should be submitted using a submission script on `cardinal`. The submission script must be run on `cardinal.stanford.edu`.
3. Copy your submission files to `cardinal.stanford.edu`. You can use the following command in your terminal:

```
scp <your submission file(s)> <your SUNetID>@cardinal.stanford.edu:
```

The script will copy the files below to a directory accessible to the CME 213 staff. Only these files will be copied. Make sure these files exist and that no other files other than those provided in the starter code are required to compile and run your code. In particular, do not use external libraries, additional header files etc, that would prevent the teaching staff from compiling the code successfully. Here is the list of files we are expecting and that will be copied:

`gpuStencil.cu`

The script will fail if one of these files does not exist. Make sure to upload `hw3.pdf` to Gradescope that is **not** the original pdf that comes with the homework!

4. Make sure your code compiles on `cme213-cluster` and runs. To check your code, we will run:
`$ make`
This should produce one executable: `main`
5. Type:
`/usr/bin/python /usr/class/cme213/WWW/script/submit.py hw3 <directory with your submission files>`
6. You can submit at most 10 times before the deadline; each submission will replace the previous one.
7. There will be a 10% penalty per 24 hours for late submission. We will not accept submissions that are submitted past two days.

B Compiling a CUDA program

B.1 Machine

We will be using the `cme213-cluster` cluster, which you can access with ssh:

```
ssh sunetid@cme213-cluster.stanford.edu
```

B.2 Compiling

To use `nvcc` on the `cme213-cluster`, you must have the correct modules loaded. If you type the command

```
$ module list
```

the output should be:

Currently Loaded Modules:

```
1) gnu/5.4.0 2) cuda/8.0 3) openmpi/1.10.6
```

If not, unload and load the appropriate modules using for example

```
$ module unload gnu/7.3.0
```

```
$ module load gnu/5.4.0
```

You can copy and paste these two lines to your `~/.bashrc` file so they will be loaded automatically when you connect on the cluster. Other useful commands include:

```
$ module avail
```

```
$ module -h
```

B.3 Running a CUDA program on the cluster

For our purposes, it will be sufficient to use: `sbatch` to submit jobs, `scancel` to delete jobs, and `squeue` to check the status. To run a job with `sbatch`, you must provide a script containing the commands you wish to run. We have provided the script `hw3.sh` for you to run your code on `cme213-cluster`. You should change the `WORKDIR` in the script. The command to submit the job script is

```
$ sbatch hw3.sh
```

You can delete your job with

```
$ scancel jobID
```

C Advice and hints

- Make sure you understand what are the parameters of the problem. In particular the difference between `nx` and `gx` is important to understand.
- Spend some time looking at the `simParams` class as it contains the useful parameters to solve the problem.
- Make sure your implementations are able to deal with square grids as well as rectangular ones.